


# <Cholesky\_kernel>

- Andrea Tirelli
- Amanjot Singh
- Danilo Bellocchio
- Domenico Rossini

# CHOLESKY DECOMPOSITION (or Cholesky Factorization)

```
for (i = 0; i < _PB_N; ++i)
{
    x = A[i][i];
    for (j = 0; j <= i - 1; ++j)
        x = x - A[i][j] * A[i][j];
    p[i] = 1.0 / sqrt(x);
    for (j = i + 1; j < _PB_N; ++j)
    {
        x = A[i][j];
        for (k = 0; k <= i - 1; ++k)
            x = x - A[j][k] * A[i][k];
        A[j][i] = x * p[i];
    }
}
```

es.



Example of data dependencies of a single column on a 5x5 matrix. Green elements are the active elements for each iteration, which directly depend on the yellow elements, indirectly on the pink ones, and not dependent of the blue ones.

# Sequential Benchmarking

Performance counter stats for './cholesky\_acc':

361.898.961	cache-references			
24.269.269	cache-misses	#	6,706 % of all	cache refs
922.929.941	cycles			
1.087.772.189	instructions	#	1,18	insn per cycle
<not supported>	branches			
581	faults			
0	migrations			

0,629194695 seconds time elapsed

Possible solution:

- **Vectorize** the problem
- Manipulate the memory **init** and **access**

# PARALLELIZED VERSION 1

```
static void kernel_cholesky(int n,  
    DATA_TYPE POLYBENCH_1D(p, N, n),  
    DATA_TYPE POLYBENCH_2D(A, N, N, n, n))  
{  
    int i, j, k;  
  
    DATA_TYPE x;  
    for (i = 0; i < _PB_N; ++i)  
    {  
        x = A[i][i];  
  
        #pragma omp parallel for reduction(-:x)  
        for (j = 0; j <= i - 1; ++j)  
            x = x - A[i][j] * A[i][j];  
  
        p[i] = 1.0 / sqrt(x);  
  
        #pragma omp parallel for private(x, k)  
        for (j = i + 1; j < _PB_N; ++j)  
        {  
            x = A[i][j];  
  
            #pragma omp simd reduction(-:x)  
            for (k = 0; k <= i - 1; ++k)  
                x = x - A[j][k] * A[i][k];  
  
            A[j][i] = x * p[i];  
        }  
    }  
}
```

Exec time: ~14s

# PARALLELIZED VERSION 2

```
static void opt_kernel_cholesky(int n,
                                DATA_TYPE POLYBENCH_1D(p, N, n),
                                DATA_TYPE POLYBENCH_2D(A, N, N, n, n))
{
    int i, j, k;

    DATA_TYPE x;
    for (i = 0; i < _PB_N; ++i) {
        p[i] = 1 / sqrt(A[i][i] - p[i]);

        #pragma omp parallel for private(j, k, x)
        for (j = i + 1; j < _PB_N; j++) {
            x = A[i][j];

            #pragma omp simd reduction(-:x)
            for (k = 0; k <= i - 1; ++k)
                x = x - A[j][k] * A[i][k];

            A[j][i] = x * p[i];
            p[j] += A[j][i] * A[j][i];
        }
    }
}
```

Exec time: ~13s

# AMDAHL'S LOW COMPARISON

## Benchmark :

Dataset = EXTRALARGE

Sequential execution time = 32 sec

---

## Amdahl's Law :

$$speedup = \frac{1}{(1-p) + \frac{p}{n}}$$

Amdahl's Law -> n = 4 p = 90%

Expected SpeedUp = 3x

---

## Parallelized Version :

Parallel Execution time = 13 sec

**Achieved SpeedUp = 2,46x**



# CONCLUSIONS

Let's analyze the results obtained in detail :

```
@@@@@ => EXTRALARGE_DATASET SEQUENTIAL
Performance counter stats for './cholesky_acc':
```

21.396.759.016	cache-references
1.799.012.184	cache-misses
47.755.147.510	cycles
64.213.298.905	instructions
<not supported>	branches
	faults
0	migrations

32,293423318 seconds time elapsed

```
@@@@@ => EXTRALARGE_DATASET PARALLEL_OPT
Performance counter stats for './cholesky_acc':
```

11.036.756.158	cache-references
1.917.290.326	cache-misses
77.290.430.309	cycles
38.568.275.253	instructions
<not supported>	branches
672	faults
31	migrations

13,145241357 seconds time elapsed

